

Neural Network Basics

José R. Dorronsoro
Dpto. de Ingeniería Informática
Escuela Politécnica Superior
Universidad Autónoma de Madrid

Modeling Basics

1 Machine Learning Basics

Basic Modeling

Regression Basics

Bias, Variance and Cross Validation

Basic Classification

Logistic Regression

2 Multilayer Perceptrons

Rosenblatt's Perceptrons

Classical MLPs

Neural Network Classification

MLP Regularization

Unconstrained Smooth Optimization

Computational Costs of MLPs

What Is Machine Learning (ML)?

- Lofty definition: make machines learn!!!
 - Have to make “machines” and “learn” more precise
- The machines of ML: mathematical input–output processes that lend themselves to some form of (numerical) parameterization
- The learning process: adjust the machine’s parameters until a goal is reached
- New thing: “goal”?
 - At first sight, get something done
 - Ultimately, to minimize some error measure
- Summing things up: a ML process tries to find a concrete mathematical/algorithmic **input–output parameterized transformation** that **minimizes an error measure** by iteratively **adjusting the transformation’s parameters**

Where Lies ML?

- In the middle of a possibly long process chain
- Before ML starts we must
 - Go from **raw to organized** data: accesing, gathering, cleaning, formatting, ...
 - Go from **organized to** (potentially) **informative** data: extracting basic and derived features
- After ML finishes and start applying a model, we must perform
 - Outcome **evaluation**: how good/actionable is
 - Outcome **exploitation**: collect, organize, act
 - **Individual model maintenance**: monitor performance, tune hyper-parameters
 - **Modeling life cycle maintenance**: discard old models, introduce new ones and **communicate** our work/results

Supervised/Unsupervised Models

- ML model types: **supervised, unsupervised**
- Supervised models:
 - Targets y^p are known and the model tries to predict or estimate them
 - These known targets guide, or **supervise**, model building
 - Main emphasis here
- Unsupervised models:
 - There are no predetermined or supervising outputs
 - But nevertheless the model is supposed to learn relations or find structure in the data
 - Often as a first step towards a supervised model
- Semisupervised models lie in between: some patterns have targets, some don't

Regression and Classification




- Problems (usually) to be solved by ML models: regression, classification
- In both patterns come in pairs (x, y)
 - x : inputs, predictors, features, independent variables
 - y : target, response, dependent variable; numerical in regression, class labels in classification
- **Regression**: the desired output y is **regressed** into the inputs x to derive a model $\hat{y} = f(x)$
 - We want $y \simeq \hat{y}$ so having $y - \hat{y}$ “small” is the natural goal
- **Classification**: inputs are derived from several classes C_1, \dots, C_K , to which labels ℓ_k are assigned
 - The model now assigns a label $\ell(x)$ to an input x
 - If x is derived from C_k we want to have $\ell(x) = \ell_k$
 - Here having $\ell(x) - \ell_k$ “small” may not make sense

The Boston Housing Regression Problem

- This is a first, widely used “toy” problem
- We want to estimate the median of house values over an area from some information about it which we believe relevant
- Features x : several real estate–related variables of Boston areas
 - CRIM: per capita crime rate by town
 - RM: average number of rooms per dwelling
 - NOX: nitric oxides concentration (parts per 10 million)
 - AGE: proportion of owner-occupied units built prior to 1940
 - LSTAT: % lower status of the population
 - ...
- Target y : MEDV, median value of owner-occupied homes in \$1,000's

The MNIST Classification Problem

- This is a well known benchmark, moderately large classification problem
- We have 28×28 grey scale image rasters with handwritten digits 0 to 9
- We want to build a classifier that assigns each raster to one of the 0 to 9 digits
- Features x : 28×28 image matrices
- Target y : labels $[0, 1, \dots, 9]$
- Examples (0-padded to 32×32):



Model Parameterization

- Often individual models are selected through (ideally optimal) **parameter sets**
 - The parameters (weights) $w \in R^M$ select a concrete f in a model family \mathcal{F}
- **Parametric** models have a fixed functional form $f(x) = f(x; W)$
 - Simplest example: **linear regression**, where x has dimension d , $M = d + 1$ and $W = (w_0, w)$

$$f(x; w_0, w) = w_0 + \sum_{j=1}^d w_j x_j = w_0 + w \cdot x$$

- **Semi-parametric** models: also use weights but without a predefined functional form
 - MLPs but also RF or GBR
- **Non parametric** models do not use weights nor follow any broad functional form
 - Nearest Neighbor models

Issues in Model Building

- There are some initial questions when working with models from a given family \mathcal{F} :
 - Which assumptions do they make?
 - How do they work on the input features?
 - How we do build them?
- In turn, these questions lead to another two:
 - How do we select the best model from the given family for the problem at hand?
 - How do we control the model building procedure?
- All of them address fundamental issues that require a **moderately deep** understanding of what is going on under the model's hood
- This understanding is usually framed in **mathematical language**

Regression Basics

1 Machine Learning Basics

- Basic Modeling

- Regression Basics**

- Bias, Variance and Cross Validation

- Basic Classification

- Logistic Regression

2 Multilayer Perceptrons

- Rosenblatt's Perceptrons

- Classical MLPs

- Neural Network Classification

- MLP Regularization

- Unconstrained Smooth Optimization

- Computational Costs of MLPs

How to Build Regression Models

- In general we have a sample $S = \{x^p, y^p\}$, $1 \leq p \leq N$, with x^p the **features** and y^p the **targets**
- We want to build a model $\hat{y} = f(x)$ so that $\hat{y}^p = f(x^p) \simeq y^p$;
 - I.e., we want to **regress** y to the $f(x)$
- The concrete f is chosen within a certain family \mathcal{F}
 - Examples here: linear regression, multilayer perceptrons (MLPs)
 - And also: SVMs, Random Forests (RF), Gradient Boosting (GB), nearest neighbor (NN)
- Natural option to ensure $f(x^p) \simeq y^p$: choose f to minimize the sample **Mean Square Error (MSE)**

$$\hat{e}(f) = \hat{e}_S(f) = \frac{1}{2N} \sum_{p=1}^N (y^p - f(x^p))^2$$

- This corresponds to the **population** error

$$\begin{aligned} e(w_0, w) &= \frac{1}{2} E_{x,y} [(y - f(x))^2] = \frac{1}{2} E[(y - f(x))^2 p(x, y) dx dy] \\ &= \frac{1}{2} \int (y - f(x))^2 p(x, y) dx dy \end{aligned}$$

Model Estimation as Error Minimization

- Thus, we may select a model by solving $\hat{f} = \hat{f}_S = \arg \min_{f \in \mathcal{F}} \hat{e}_S(f)$
 - But this may be quite complicated mathematically and we must simplify it
- For a parametric or semiparametric $f(x; W)$ we can write $\hat{e}_S(f) = \hat{e}_S(f(\cdot; W)) = \hat{e}_S(W)$
- The problem to solve becomes

$$\hat{W}^* = \hat{W}_S^* = \arg \min_W \hat{e}_S(W), \text{ i.e., } \hat{e}_S(\hat{W}^*) \leq \hat{e}_S(W) \forall W$$

- For linear regression (LR) the sample error is

$$\hat{e}(w_0, w) = \frac{1}{2N} \sum_p (y^p - w_0 - w \cdot x^p)^2$$

- Minimizing the sample error $\hat{e}_S(W)$ becomes a problem in mathematical **optimization**

Regression Assumptions

- **Key assumption:** x and y are related as $y = \phi(x) + n$ where
 - $\phi(x)$ is the **true** underlying function
 - n is **additive noise** with 0 mean and finite variance σ_N^2
- Our sample is just a particular instance of a deeper **sample generation process**
- Thus x, n are produced by **random variables** X, N
 - And so is y , given by $Y = \phi(X) + N$
- Moreover, X and N are **independent distributions** with densities $q(x), \nu(n)$
- Thus, X and Y (or X and N) have a joint density

$$p(x, y) = p(x, \phi(x) + n) = q(x) \nu(n) = q(x) \nu(y - \phi(x))$$

MSE Decomposition

- We can decompose the MSE error of any model f as

$$\begin{aligned} 2\text{mse}(f) &= E_{x,y}[(y - f(x))^2] = \int (n + \phi(x) - f(x))^2 q(x) \nu(n) dx dn \\ &= \int (n^2 + 2n(\phi(x) - f(x)) + (\phi(x) - f(x))^2) q(x) \nu(n) dx dn \\ &= \int n^2 \nu(n) dn + \int (\phi(x) - f(x))^2 q(x) dx + \\ &\quad 2 \left(\int n \nu(n) dn \right) \left(\int (\phi(x) - f(x)) q(x) dx \right) \\ &= \sigma_N^2 + E_x[(\phi(x) - f(x))^2] \end{aligned}$$

- Thus **for any model we have $\text{mse}(f) \geq \sigma_N^2$ always**
 - Beware of zero regression errors!!
- And we should focus on achieving on $f \simeq \phi$ (which we don't know!!)

The Best Regression Model


- It is easy to see that the best f is simply $f(x) = E_y[y|x]$, for

$$E_y[y|x] = E_n[\phi(x) + n] = \int (\phi(x) + n) \nu(n) dn = \phi(x)$$

- Have we finished? In theory yes; in practice, not at all!!!
 - We do not know ν and, thus, cannot compute the required integral
 - If we would have several M values y^j for any x , we could try $\hat{\phi}(x) = \frac{1}{M} \sum_1^M y^j$
 - But this doesn't happen either
- Now we have two options:
 - Try to stretch the $E[y|x]$ approach, for instance, by **k -NN Regression**
 - Forget about it and get back to get models f for which we hope $f \simeq \phi$

From The Best Regressor to k -NN

- The left hand side shows the ideal situation but the right one is what we should expect in practice



k -NN Regression

- We can try to use the y^p values of several x^p close to a new x
- We fix k neighbors x^{p_1}, \dots, x^{p_k} of x and estimate $\hat{y} = \hat{y}(x)$ as

$$\hat{y}(x) = \frac{1}{k} \sum_{j=1}^k y^{p_j}$$

- $\hat{y}(x) = \hat{Y}_k^{NN}(x)$ is the k -**Nearest Neighbor (NN)** regressor
 - This can be refined to weighted versions, such as

$$\hat{y}(x) = \frac{1}{C_k(x)} \sum_{j=1}^k \frac{1}{\|x^{p_j} - x\|^2} y^{p_j}$$

with $C_k(x) = \sum_{j=1}^k \frac{1}{\|x^{p_j} - x\|^2}$ a normalizing constant

- But we need that close predictors give predictions that are also close, and that **there are enough of them close by**
 - This is **very unlikely**

The Curse of Dimensionality

- Even for low dimensions and large samples, **the sample space is essentially empty**
- Assume we have 1,000 d -dimensional x patterns whose features have values between 1 and 10
 - In dimension $d = 1$ there are 100 patterns per unit length
 - But when $d = 3$ we have just 1 pattern per volume unit
 - And if $d = 6$ we have just 1 pattern per 1,000 units of volume
 - And in dimension 10 (not a big one nowadays) we have just ... !!!
- Thus, for most problems, **there never will be enough close points**
- As a consequence, to get k observations we may go too far away from x and the average will not be meaningful
- Therefore, unless we deal with violently non-linear problems, a simple linear model may be better than k -NN regression for moderate dimensions

Linear Models

- Assuming $x \in R^d$, recall the basic linear model is

$$f(x) = w_0 + \sum_1^d w_i x_i = w_0 + w \cdot x$$

- w_0 complicates notation; to drop it we center x and y so that $E[x_i] = E[y] = 0$; then $w_0 = 0$
- Then we are left with the simpler **homogeneous** model $f(x) = w \cdot x$
- In practice we will always **normalize** x , for instance to have 0 mean and 1 standard deviation (std) on each feature
 - But not y if we may help it (or it is easy to reverse this)
- But: how do we find w ?

1-dimensional Linear Regression (LR)

- Assume that features X and target Y are **centered**, i.e., have 0 means
- For 1-dimensional patterns x the LR model then becomes $f(x) = w x$
- And the error is then the function

$$\begin{aligned}\hat{e}(w) &= \frac{1}{2N} \sum_{p=1}^N (w x^p - y^p)^2 = \frac{1}{2N} \sum_p (w^2 (x^p)^2 - 2x^p y^p w + (y^p)^2) \\ &= \frac{1}{2} \left(\frac{1}{N} \sum_p (x^p)^2 \right) w^2 - \left(\frac{1}{N} \sum_p x^p y^p \right) w + \frac{1}{2N} \sum_p (y^p)^2\end{aligned}$$

- Thus, $\hat{e}(w) = aw^2 + bw + c$ with $a > 0$ and it has obviously a minimum w^*
- To find it we just solve $\hat{e}'(w) = 0$

Solving $\hat{e}'(w) = 0$

- We first compute $\hat{e}'(w)$, for which we have

$$\hat{e}'(w) = w \left(\frac{1}{N} \sum_p (x^p)^2 \right) - \frac{1}{N} \sum_p x^p y^p$$

- The optimal w^* solves $\hat{e}'(w) = 0$ and is given by

$$w^* = \frac{\frac{1}{N} \sum_p x^p y^p}{\frac{1}{N} \sum_p (x^p)^2} = \frac{\frac{1}{N} X^t Y}{\frac{1}{N} X^t X} = \frac{\text{covar}(x, y)}{\text{var}(x)}$$

where X and Y denote the $N \times d$ **data matrix** (vector here) and the $N \times 1$ **target vector**

General Linear Regression

- Assume again that X and Y are centered
- The LR model becomes now $f(x) = \sum_1^d w_i x_i = w \cdot x$
- If Y is the $N \times 1$ **target** vector and we organize the sample S in a $N \times d$ **data matrix** X , the sample mse is given by

$$\begin{aligned}\hat{e}(w) &= \frac{1}{2N} \sum_p (w \cdot x^p - y^p)^2 = \frac{1}{2N} (Xw - Y)^t (Xw - Y) \\ &= \frac{1}{2N} (w^t X^t X w - 2w^t X^t Y + Y^t Y)\end{aligned}$$

- We have now a **quadratic form** $w^t A w + w^t b + c$, with A a $d \times d$ matrix, b a d vector and c a scalar
- Now we have to solve $\nabla \hat{e}(w) = 0$, i.e., $\frac{\partial \hat{e}}{\partial w_i}(w) = 0, 1 \leq i \leq d$
- It is easy to see that

$$\nabla \hat{e}(w) = \frac{1}{N} X^t X w - \frac{1}{N} X^t Y = \hat{R} w - \hat{b}$$

Solving the Linear Equations

- The optimal \hat{w}^* must verify $\nabla \hat{e}(\hat{w}) = \hat{R} \hat{w} - \hat{b} = 0$, where

$$\hat{R} = \frac{1}{N} X^t X, \quad \hat{b} = \frac{1}{N} X^t Y$$

- Over the original, non-centered data matrix we have

$$\hat{R} = \frac{1}{N} (X - \bar{X})^t (X - \bar{X});$$

i.e., \hat{R} is the **sample covariance matrix**

- If \hat{R} is **invertible**, we just solve the linear system $\hat{R} \hat{w} = \hat{b}$ to get

$$\hat{w}^* = \hat{R}^{-1} \hat{b} = (X^t X)^{-1} X^t Y = \text{covar}(X)^{-1} \text{covar}(X, Y)$$

- We won't but we should write \hat{w}_S^* , as it is **sample-dependent**
- \hat{R} is oftent inverted through the **Singular Value Decomposition** (SVD) of the data matrix X

SVD

- The **Singular Value Decomposition** (SVD) of a (complex) $m \times n$ matrix M , $m \geq n$, is $M = UDV^*$ where
 - U and V are orthonormal (unitary) matrices with dimensions $m \times m$, $n \times n$ respectively
 - D is an $m \times n$ matrix, diagonal on its first n rows and with zeros in its last $m - n$ ones
- Graphically, the SVD decomposition of M is

$$\begin{matrix} \text{4x4} & & \text{4x4} & & \text{4x4} & & \text{4x4} \\ \mathbf{M} & = & \mathbf{U} & \mathbf{\Sigma} & \mathbf{V}^* \\ m \times n & & m \times m & & m \times n & & n \times n \end{matrix}$$

Taken from *Wikipedia*

- The columns of U and V are called the **left and right singular vectors**; the diagonal elements in D are the **singular values**
- We may interpret this image as M performing first a rotation with V^* , then feature dilations with Σ (our D) and, finally, another rotation with U
 - Actually both are a roto-reflections

SVD II

- It is easy to see that we can also write $M = \tilde{U}\tilde{D}V^*$, with \tilde{U} having the first n columns of U and \tilde{D} the top diagonal block of D
 - Assuming M to be real and with a slight abuse of language, we will write $M = UDV^t$
- In our case we can write $X = UDV^t$, with dimensions $N \times d$, $d \times d$ and $d \times d$ respectively
- Then $X^tX = VDU^tUDV^t = VD^2V^t$ and $(X^tX)^{-1} = VD^{-2}V^t$
 - D^{-2} is a diagonal matrix with $(D^{-2})_{ii} = 1/D_{ii}^2 = 1/\delta_i^2$
- The columns of U and V are unique up to a sign change
 - So different packages may give seemingly different U and V

Gradient Descent MSE Minimization

- Computing the covariance matrix has a $O(N \times d^2)$ cost and invert it has a $O(d^3)$ cost
 - For big data problems it may not be possible to solve analytically the normal equation $\nabla \hat{e}(w) = 0$

- The simplest numerical alternative is **gradient descent**:

- Starting from some random w^0 , we iteratively compute

$$w^{k+1} = w^k - \rho_k \nabla \hat{e}_B(w^k) = w^k - \frac{\rho}{n_B} (X_B^t X_B w^k - X_B^t Y)$$

over a **mini-batch** X_B with n_B samples and batch error \hat{e}_B

- Component wise: $w_i^{k+1} = w_i^k - \rho_k \frac{\partial \hat{e}_B}{\partial w_i}(w^k)$
- ρ_k is the **learning rate**
- If $w^k \rightarrow w^*$, then $\nabla \hat{e}(w^*) = 0$
 - Since our problems have obviously minima, this should be enough

By the Way ...

- We have just built our first neural network: the linear regression architecture can be seen as a multilayer perceptron (MLP) with no hidden layers
- Thinking of this as an MLP seems far-fetched, as we are missing the hidden layers and the much more complicated gradient computation
- But the layer architecture is there and we also have learning as iterative error minimization, i.e., **training**
- In fact, we will see that MLPs perform **linear regression on the last hidden layer outputs**
- And the process from the input to the last hidden layer produces an **enhanced feature representation** upon which the regression error should be better

Measuring Model Fit

- First option: **Root Square Error** $RSE = \sqrt{\frac{1}{N} \sum (y^p - \hat{y}^p)^2}$
- OK, but how to compare this? We must always have a **base model** to benchmark our results
- Simplest model: a constant $\hat{y} = w_0$, which reduces to predict the mean $\bar{y} = \frac{1}{N} \sum_1^N y^p$, with square error

$$\frac{1}{N} \sum (y^p - \bar{y})^2 = \text{Var}(y)$$

- We can compare our model against this base model by computing

$$\frac{\sum (y^p - \hat{y}^p)^2}{\sum (y^p - \bar{y})^2} = \frac{\frac{1}{N} \sum (y^p - \hat{y}^p)^2}{\frac{1}{N} \sum (y^p - \bar{y})^2} = \frac{RSE^2}{\text{Var}(y)}$$

- The widely used R^2 coefficient is simply $R^2 = 1 - \frac{RSE^2}{\text{Var}(y)}$

Regularization

- Our regression solution $\hat{w}^* = (X^t X)^{-1} X^t Y$ won't work if $X^t X$ is not invertible
 - For instance, when some features are highly correlated
 - We can at most say that $X^t X$ is **positive semidefinite**
- We could fix this working instead with $X^t X + \alpha I$ for some $\alpha > 0$
 - This is then always **positive definite** and we can thus invert it
- To make this practical, note that $\hat{w}^* = (X^t X + \alpha I)^{-1} X^t Y$ minimizes

$$e_R(w) = \frac{1}{2N} \sum_p (y^p - w \cdot x_p^p)^2 + \frac{\alpha}{2} \|w\|^2,$$

- This is the **Ridge Regression** problem
 - Our first example of **regularization**, a key technique in Machine Learning
 - **All ML models must be regularized in some way**
- Important issue: how to find the right choice for α ?

Takeaways on Linear Regression

- ① We introduced **supervised** models
- ② We have reviewed the essentials of the **linear regression model** (always the first thing to try)
- ③ We have considered model estimation as a problem on **error minimization**
- ④ We have seen how to build linear models **analytically and numerically**
- ⑤ We have seen how to **measure model fit**
- ⑥ We have introduced **regularization**

Bias, Variance and Cross Validation

1 Machine Learning Basics

Basic Modeling

Regression Basics

Bias, Variance and Cross Validation

Basic Classification

Logistic Regression

2 Multilayer Perceptrons

Rosenblatt's Perceptrons

Classical MLPs

Neural Network Classification

MLP Regularization

Unconstrained Smooth Optimization

Computational Costs of MLPs

Sample Bias and Variance

- Remember: all the models f we will build are **sample dependent**, i.e., $f = f_S$
- With several **independent** samples S_1, \dots, S_M , it is natural to use as our best final model the average of their associated $\hat{f}_{S_m}(x)$ models, i.e.,

$$\frac{1}{M} \sum_1^M \hat{f}_{S_m}(x) \simeq E_S[\hat{f}_S(x)] = \hat{f}_N(x)$$

- Here we assume x fixed and take the (theoretical) expectation $E_S[\hat{f}_S(x)]$ over all samples S of size N
- $\hat{f}_N(x) = E_S[\hat{f}_S(x)]$ is our ideal **best model** and $\phi(x) - \hat{f}_N(x)$ is its **bias**
- The **variance** of the $\hat{f}_S(x)$ estimates is then

$$V_N(x) = E_S [(\hat{f}_S(x) - \hat{f}_N(x))^2]$$

Bias Versus Variance

- Recall that our goal is to estimate ϕ in our regression model $y = \phi(x) + n$
- Ideally we would like to have a model such that


$$\hat{f}_N(x) - \phi(x) \simeq 0,$$

i.e., a model with small **bias**

- This should be achievable if we have rich, highly flexible models
 - Or with essentially no regularization (or both!)
- But we would also like to have a robust model building procedure which results in a small **variance** $V_N(x) \simeq 0$
 - This should be achievable if models are simple with few parameters
 - Or with more severe regularization if not
- Obviously both goals are contradictory to a large extent

The Bias–Variance Tradeoff

- There is thus a **tradeoff** between bias (low for complex models) and variance (low for simple models)



Taken from *Hastie et al.*, p. 38

Bias vs Variance in k -NN Regression

- Recall that in k -NN regression we choose the k nearest neighbors x^{p_1}, \dots, x^{p_k} of x and estimate $\hat{y} = \hat{y}(x)$ as

$$\hat{y}(x) = \frac{1}{k} \sum_{j=1}^k y^{p_j}$$

- Here the parameter that controls the tradeoff is just k
- If $k = N$, the sample size, the N -NN estimator is just the mean: $Y_N^{NN} = \bar{y}$, with very small variance but large bias (it's an obviously bad model!!)
- If $k = 1$, the 1-NN estimator will have smaller bias but a large variance: changing the sample is very likely to change the sample point nearest to x

Bias vs Variance in Ridge Regression

- Recall that the Ridge Regression loss is

$$e_R(w) = \frac{1}{2N} \sum_p (y^p - w \cdot x_p^p)^2 + \frac{\alpha}{2} \|w\|^2,$$

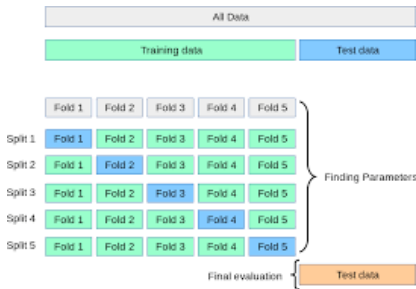
- Here the parameter that controls the tradeoff is α
- If $\alpha \gg 1$, any non zero w implies a large regularization penalty
 - It is thus likely that $w \simeq 0$ and the Ridge model reduces again to the mean \bar{y} , with small variance but possibly a large bias
- But if $\alpha \simeq 0$, w can wander on the entire \mathbf{R}^d
 - The bias will be then smaller, but the weights w_S and $w_{S'}$ from different samples may be very different, particularly if the covariance is nearly singular

Evaluating Expected Performance

- Recall that we build a model over a **train** sample but apply it on a **test** sample
- Thus, before we start applying a model, we should have a reasonably accurate idea of its performance in practice
- I.e., we want to estimate the model's **generalization performance**
- Estimating the generalization performance **only over the sample S used for training results in misleading error values**
- The preceding suggests to try to have M independent subsamples S_m and then
 - To compute $\hat{f}_M(x) = \frac{1}{M} \sum_m \hat{f}_{S_m}(x) \simeq \hat{f}_N(x)$
 - To get the error estimate $\hat{e} = \frac{1}{N} \sum_p (y^p - \hat{f}_M(x^p))^2$ over a new, **unseen** sample $S' = \{(x^p, y^p)\}$
- But since usually we only have a single S , we “simulate” this by **Cross Validation** (CV) to get our first realistic generalization error estimates

Cross Validation I

- Graphically, the general schema of Cross Validation is the following



From [Scikit-learn](#)

Cross Validation II

- In more detail, in Cross Validation (CV) and for a global sample S we
 - Randomly split $S = S_{tr-val} \cup S_{ts}$
 - Randomly split S_{tr-val} in M subsets S_1, \dots, S_M
 - Work with M **folds**: pairs (S_m, S_m^c) , with

$$S_m^c = S - S_m = \cup_{i \neq m} S_i$$

- Build M different models **using the S_m^c as training subsets**
 - Compute their errors e_m on the folds' **validation subsets S_m**
 - Use these errors' average as a first estimate of the true model performance
 - Test it if possible on S_{ts}
- CV can and **must be used** in any model building procedure
 - Most data science packages have tools to simplify this
- We will also use CV to find **optimal model hyper-parameters** such as α in Ridge Regression

Grid Hyper-parameter Selection

- Consider for Ridge regression a hyperparameter range $\alpha \in [a, A]$, with $0 < a < A$
 - In general $a \simeq 0$, $A \gg 1$
 - $\alpha = a$: essentially no penalty, small bias and high variance
 - $\alpha = A$: large penalty, small variance but high bias
- Select an $L + 1$ point **grid** $\{\alpha_0 = a < \alpha_1 < \dots < \alpha_L = A\}$
 - The α_j can be equi-spaced, log equi-spaced, random, ...
- At each α_ℓ
 - Train M models on the S_m^c using the hyperparameter α_ℓ
 - Average their test errors e_m on the S_m to get the error $e(\alpha_\ell)$ at α_ℓ
- And choose the final (hopefully) optimal hyperparameter α^* as

$$\alpha^* = \arg \min_{0 \leq \ell \leq L} e(\alpha_\ell)$$

- α^* gives the model with the **best expected generalization among all possible α choices**

Takeaways on Bias, Variance and CV

- ① We have stressed that **any model estimation is sample-dependent** and that this has to be controlled
- ② We have introduced the **bias** and **variance** as the two key components of any model error
- ③ We have discussed **bias-variance trade-off**
- ④ We have introduced **Cross Validation** here as a tool to estimate a **model's generalization performance**
- ⑤ We have also introduced **Cross Validation** as a tool to estimate a **model's hyper-parameters**

The Classification Model

1 Machine Learning Basics

Basic Modeling

Regression Basics

Bias, Variance and Cross Validation

Basic Classification

Logistic Regression

2 Multilayer Perceptrons

Rosenblatt's Perceptrons

Classical MLPs

Neural Network Classification

MLP Regularization

Unconstrained Smooth Optimization

Computational Costs of MLPs

Regression vs Classification

- Recall that in regression we have numerical continuous targets y and want our predictions \hat{y} to be as close to y as possible
- But in classification we have a finite number of labelled targets for which “selection by closeness” doesn’t make sense
- Natural alternative: select the **most probable** label given the pattern x we have just received
 - The concrete labels used for targets do not matter much anymore
 - Model learning should thus be “target” agnostic
 - And good probability estimates should be quite useful
- Let’s analyze this in an example

A First Problem: Pima Indian Diabetes

- We want to diagnose whether a person may have diabetes from some clinical measures
- Features x : clinical measures
 - numPregnant
 - bloodPress
 - massIndex
 - age ...
- Target y : 0 (no diabetes), 1 (diabetes)
- Clear goal but perhaps too radical
- Better: try to estimate the probability $P(1|x)$ of having diabetes depending on the features x we measure

Classification Setup

- We have random patterns ω from M classes, C_1, \dots, C_M
- Over each pattern we “measure” d features $x = x(\omega) \in \mathbb{R}^d$
 - x inherits the randomness in ω and becomes a random variable
- A ω has a **prior probability** π_m of belonging to C_m
- Inside each class C_m there is a **conditional class density** $f(x|m)$ that “controls” the appearance of a given x
- The π_m and $f(x|m)$ determine the **posterior probability** $P(m|x)$ that x comes from class C_m
- **Intuition:** we should assign x to the class with the largest $P(m|x)$, that is, work with the classifier

$$\delta(x) = \arg \max_m P(m|x)$$

Computing Posterior Probabilities I

- **Bayes rule:** $P(B|A) = \frac{P(A \cap B)}{P(A)}$
- This requires to work with probabilities, not densities, but $P(\{x\}) = P(m \cap \{x\}) = 0$ and

$$P(m|x) = \frac{P(m \cap \{x\})}{P(\{x\})} = \frac{0}{0} = \dots???$$

- But we can use the approximation

$$\begin{aligned} P(m|x) &\simeq P(m|B_r(x)) = \frac{P(C_m \cap B_r(x))}{P(B_r(x))} = \frac{P(B_r(x)|m)P(C_m)}{P(B_r(x))} \\ &= \frac{\pi_m P(B_r(x)|m)}{P(B_r(x))} = \pi_m \frac{\int_{B_r(x)} f(y|m) dy}{\int_{B_r(x)} f(z) dz} \end{aligned}$$

where we assume that features x are measured independently from classes m

Computing Posterior Probabilities II

- Remember the Fundamental Theorem of Calculus:
if $F(x) = \int_a^x f(y)dy$,

$$\lim_{\epsilon \rightarrow 0} \frac{1}{2\epsilon} \int_{x_0-\epsilon}^{x_0+\epsilon} f(y)dy = \frac{dF}{dx}(x_0) = f(x_0)$$

- In d dimensions it becomes

$$g(w) = \lim_{r \rightarrow 0} \frac{1}{|B_r(w)|} \int_{B_r(w)} g(z)dz$$

- Putting everything together, we arrive

$$\begin{aligned} P(m|x) &= \lim_{r \rightarrow 0} P(m|B_r(x)) = \pi_m \lim_{r \rightarrow 0} \frac{\int_{B_r(x)} f(y|m)dy}{\int_{B_r(x)} f(z)dz} \\ &= \pi_m \lim_{r \rightarrow 0} \frac{\frac{1}{|B_r(x)|} \int_{B_r(x)} f(y|m)dy}{\frac{1}{|B_r(x)|} \int_{B_r(x)} f(z)dz} = \frac{\pi_m f(x|m)}{f(x)} \end{aligned}$$

The Obviously Optimal Classifier

- This suggests that a natural choice is the **classifier** δ_B

$$\begin{aligned}\delta_B(x) &= \arg \max_m P(m|x) = \arg \max_m \frac{\pi_m f(x|m)}{f(x)} \\ &= \arg \max_m \pi_m f(x|m)\end{aligned}$$

- With some extra work we can show that this **Bayes Classifier** δ_B defines an optimal solution (in some precise sense) of the classification problem
- But ... this doesn't look too practical, for we do not know either π_m or (much harder) $f(x|m)$

Approximating the Bayes Classifier

- To define δ_B we need to know the prior probabilities π_m and the prior densities $f(x|m)$
- A reasonable choice for π_m is $\hat{\pi}_m = \frac{N_m}{N}$, where N_m is the number of patterns of C_m in the sample
- But effective multidimensional density estimates are rather difficult, because of the **curse of dimensionality**
 - Densities generalize histograms
 - Good histograms need accurate counts of elements nearby
 - But in high dimensions there won't be nearby elements!!
- Options:
 - Restrict possible density models: **logistic regression**
 - Assume no model and apply a **Nearest Neighbor (NN) strategy**

The k -NN Classifier

- Very simple: at any x consider the subset $N_k(x)$ of its k closest sample points and
 - Let $n_m(x)$ the number of elements of class m in $N_k(x)$
 - Notice that $0 \leq n_m(x) \leq k$
 - Define $\delta_{kNN}(x) = \arg \max_m n_m(x)$
- That is, $\delta_{kNN}(x)$ assigns x to the class that has more patterns in $N_k(x)$
- We can partially justify this definition from a Bayesian point of view
- Assume that $B_r(x)$ is the smallest ball that contains $N_k(x)$ and consider the approximations
 - $P(B_r(x)|m) \simeq \frac{n_m(x)}{N_m}$
 - Similarly, $P(B_r(x)) \simeq \frac{k}{N}$
 - And $\pi_m \simeq \frac{N_m}{N}$

A k -NN Example

- Assume $N = 200$, $N_+ = 150$, $N_o = 50$
- Set $k = 5$, fix x and assume $n_+(x) = 3$, $n_o(x) = 2$
- Then $\pi_+ = \frac{3}{4}$, and for our hypothetical $B_r(x)$ we have
 - $P(B_r(x)) = \frac{5}{200} = \frac{1}{40}$
 - $P(B_r(x)|+) = \frac{3}{150} = \frac{1}{50}$
- And, therefore

$$\begin{aligned} P(+|x) &\simeq P(+|B_r(x)) = \frac{\pi_+ P(B_r(x)|+)}{P(B_r(x))} = \frac{\frac{3}{4} \frac{1}{50}}{\frac{1}{40}} \\ &= \frac{3}{5} \end{aligned}$$

k -NN and the Bayes Classifier

- We then have

$$\begin{aligned} P(m|x) &\simeq P(m|B_r(x)) = \frac{\pi_m P(B_r(x)|m)}{P(B_r(x))} \\ &\simeq \frac{N_m}{N} \frac{n_m(x)}{N_m} \frac{1}{\frac{k}{N}} = \frac{n_m(x)}{k} \end{aligned}$$

- Therefore δ_{kNN} might be close to δ_B , for

$$\begin{aligned} \delta_{kNN}(x) &= \arg \max_m n_m(x) = \arg \max_m \frac{n_m(x)}{k} \\ &\simeq \arg \max_m P(m|x) = \delta_B(x) \end{aligned}$$

Some k -NN Issues

- **Q1: How do we choose k ?** Using CV, of course
- There are no closed form solution and we have to deal again with the bias–variance tradeoff
 - Small variance with large k : if $k = N$, k -NN classification returns the majority class
 - Small bias with small k : if $k = 1$ a point very close to x should be in the same class
 - But also large variance: the nearest point to x in another sample may well belong to a different class
- **Q2: Is k -NN always meaningful?**
 - Well, only if classes are concentrated and far from each other
 - In other words, when the classification problem is easy!!

Measuring Classifier Accuracy

- Consider a two class problem with labels $y = 0, 1$
- We will call patterns with label 1 **positive** and those with label 0 **negative**
 - Usually the positive patterns are the interesting ones: sick people, defaulted loans, ...
- Let $\hat{y} = \hat{y}(x)$ the label predicted at x ; we say that x is a
 - **True Positive** (TP) if $y = \hat{y} = 1$
 - **True Negative** (TN) if $y = \hat{y} = 0$
 - **False Positive** (FP) if $y = 0$ but $\hat{y} = 1$
 - **False Negative** (FN) if $y = 1$ but $\hat{y} = 0$
- These values are combined in many performance measures

Classifier Metrics

- The classifier **accuracy** is $acc = \frac{TP+TN}{N}$
 - The first thing to measure but ...
- It may not be too significant when the number N_0 of negatives is $\gg N_1$, the number of positives; then
 - The classifier $\delta(x) = 0$ will have a high accuracy $N_0/N \simeq 1$
 - But it will also be useless!!
- First variant: Precision, Recall
 - **Recall**: $TP/(TP + FN)$, i.e., the fraction of positives detected
 - **Precision**: $TP/(TP + FP)$, i.e., the fraction of true alarms issued
- Recall measures how many positive cases we recover, i.e., how effective is our method
- Precision measures the effort we need for that, i.e., its efficiency
- Ideally: high recall, high precision (i.e., effective and efficient!!)
- Another way of presenting the performance is through the **confusion matrix**

The Confusion Matrix

- Standard layout

	P' (Predicted)	N' (Predicted)
P (Actual)	True Positive	False Negative
N (Actual)	False Positive	True Negative

- Other layouts:
 - **Positives (with label 1) at bottom** (as done in `confusion_matrix` of `sklearn`)
 - Predicted values in rows, real values in columns
- Accuracy and the confusion matrix can be extended to multi-class problems
 - But not other concepts such as recall or precision

Logistic Regression (LogR)

1 Machine Learning Basics

Basic Modeling

Regression Basics

Bias, Variance and Cross Validation

Basic Classification

Logistic Regression

2 Multilayer Perceptrons

Rosenblatt's Perceptrons

Classical MLPs

Neural Network Classification

MLP Regularization

Unconstrained Smooth Optimization

Computational Costs of MLPs

Linear Regression for Classification?

- k -NN Classifier is simple but also crude; have to look elsewhere
- Building a regression model with targets given by some coding of class labels usually doesn't make sense
- However, for a binary 0–1 response, it can be shown that the $w_0 + w \cdot x$ obtained using linear regression is in fact an estimate of $P(1|x)$
 - We may thus fix a threshold δ_0 and decide 0 if $w_0 + w \cdot x < \delta_0$ and 1 otherwise
 - However, we may end up with probability estimates less than 0 or bigger than 1!!!
- Better idea: try to transform the linear output $w_0 + w \cdot x \in (-\infty, \infty)$ into a probability $P(1, x) \in (0, 1)$
 - Logistic regression does this through the **logistic** or **sigmoid** function

Logistic Regression (LogR)

- We assume

$$P(1|x) = P(1|x; w_0, w) = \frac{1}{1 + e^{-(w_0 + w \cdot x)}}$$

- Then $0 \leq P(1|x) \leq 1$ for any x
- We then have

$$P(0|x) = 1 - P(1|x) = \frac{e^{-(w_0 + w \cdot x)}}{1 + e^{-(w_0 + w \cdot x)}} = \frac{1}{1 + e^{w_0 + w \cdot x}}$$

- Notice that if $w_0 + w \cdot x = 0$, $P(1|x) = P(0|x) = 0.5$
- The ratio $\frac{P(1|x)}{P(0|x)} = e^{w_0 + w \cdot x}$ is called the **odds** of x and its log the **log odds** or **logit**
- Thus, the basic assumption in LogR is that the **logit is a linear function** $w_0 + w \cdot x$ of x
- We have the model $f(x; w)$; we need a **loss** function $L(w_0, w)$ to minimize for which we use the sample's **likelihood**

Estimating w_0^* , w^*

- Assume a single sample x, y and two possible model coefficients w_0, w and w'_0, w'
- Denoting by $p = P(y|x; w_0, w)$ and $p' = P(y|x; w'_0, w')$, it is clear that we should prefer w_0, w if $p > p'$ and w'_0, w' if not
 - In other words, we prefer the coefficients that give a **higher posterior probability**
- For an **independent** sample $S = \{(x^p, y^p)\}$, its joint probability under a posterior model $p = P(y|x, w_0, w)$ is

$$P(Y|X; w_0, w) = \prod_{p=1}^N P(y^p|x^p; w_0, w)$$

- And, again, given two possible model coefficients w_0, w and w'_0, w' , **we should prefer w_0, w iff**

$$P(Y|X; w_0, w) > P(Y|X; w'_0, w')$$

Sample's Likelihood

- Therefore, we can estimate the optimal w_0^* , w^* as

$$w_0^*, w^* = \arg \max_{w_0, w} P(Y|X; w_0, w)$$

- By the independence assumption we have

$$\begin{aligned} P(Y|X; w_0, w) &= \prod_{p=1}^N P(y^p|x^p; w_0, w) \\ &= \left\{ \prod_{y^p=1} P(1|x^p) \right\} \left\{ \prod_{y^p=0} P(0|x^p) \right\} \\ &= \prod_{p=1}^N P(1|x^p)^{y^p} P(0|x^p)^{1-y^p} \end{aligned}$$

because

- If $y^p = 1$, $P(1|x^p) = P(1|x^p)^{y^p} = P(1|x^p)^{y^p} P(0|x^p)^{1-y^p}$, and
- If $y^p = 0$, $P(0|x^p) = P(0|x^p)^{1-y^p} = P(1|x^p)^{y^p} P(0|x^p)^{1-y^p}$

Max Log-Likelihood Estimation

- The **log-likelihood** of w_0, w given S is then

$$\begin{aligned}\ell(w_0, w; S) &= \log P(Y|X; w_0, w) \\ &= \sum_p \{y^p \log p(1|x^p) + (1 - y^p) \log p(0|x^p)\} \\ &= \sum_p y^p \log \frac{p(1|x^p)}{p(0|x^p)} + \sum_p \log p(0|x^p) \\ &= \sum_p y^p (w_0 + w \cdot x^p) - \sum_p \log(1 + e^{w_0 + w \cdot x^p})\end{aligned}$$

- We can thus estimate the optimal \hat{w}_0^*, \hat{w}^* as

$$\hat{w}_0^*, \hat{w}^* = \arg \min_{w_0, w} -\ell(w_0, w; S)$$

- Extra bonus: $-\ell$ is a convex differentiable function of (w_0, w) and, thus, it is enough to solve $\nabla \ell(w_0, w) = 0$

Newton–Raphson Solution

- However, $\nabla \ell(\mathbf{w}_0, \mathbf{w}) = 0$ doesn't admit a closed form solution, but only an iterative, numerical one
- We apply the **Newton–Raphson** iterative method, here equivalent to the general Newton method for function minimization
- For a general weight vector \mathbf{w} , and starting with an initial random \mathbf{w}^0 , Newton's iterations are

$$\mathbf{w}^{k+1} = \mathbf{w}^k + (\mathcal{H}_\ell(\mathbf{w}^k))^{-1} \nabla \ell(\mathbf{w}^k)$$

- $\mathcal{H}_\ell(\mathbf{w}^k)$ denotes the **Hessian** of ℓ at \mathbf{w}^k , which may or may not be invertible
 - Everything is fine if the \mathbf{w}^k are close enough to the optimum \mathbf{w}^* but far away things may get tricky
- Just as before, we can add a regularization term $\frac{\alpha}{2} \|\mathbf{w}\|^2$ to avoid invertibility problems
- The iterations in Logistic Regression are again typical of many of the model building methods used in Machine Learning

Recap: Learning in ML

- The general approach to **learning** is the following:
 - A **model** $f(x; w)$ is chosen
 - Given a sample $S = \{(x^1, y^1), \dots, (x^N, y^N)\}$, we define a **sample dependent loss function**

$$L(w) = L(w|S) = L(y^1, \dots, y^N, f(x^1; w), \dots, f(x^N; w))$$

- $L(w)$ is often minimized from some w^0 by **iterations**

$$w^{k+1} = w^k - \rho_k G(w^k, S)$$

with ρ_k a **learning rate** and G some vectorial function

- When $G(w) = \nabla L(w)$ we have **gradient descent**
 - When $G(w) = \mathcal{H}(w)^{-1} \nabla L(w)$ we obtain **Newton's method**
- When the entire sample S is used at each iteration, we speak of **batch learning**
- When only single patterns (x^p, y^p) or small subsamples are used, we speak of **on-line** or **minibatch learning**

By the Way ...

- We have just built our second neural network: the logistic regression architecture can be seen as a multilayer perceptron (MLP) with no hidden layers and sigmoid outputs
- Again, thinking of this as an MLP seems far-fetched, as we are missing the hidden layers
- Now the gradient computations are more involved
- But we also have learning as iterative **loss** minimization, i.e., **training**
- In fact, we will see that classification MLPs for two class problems perform **logistic regression on the last hidden layer outputs**
- And , again, the process from the input to the last hidden layer produces an **enhanced feature representation** upon which the regression error should be better

What's New from Regression?

- Some things change from regression, some don't
- We should check feature correlations: although less clearly influential, they will affect most models
 - Scatter plots (x_i, y) are usually less informative
- Important: **positive and negative-class feature histograms**
- The **bias-variance trade-off** is subtler in classification
- Accuracy, recall, precision are the usual model quality measures
 - But NN training does not directly minimize them
- We use CV with **stratified folds** to estimate generalization performance
- We also use CV for hyperparameter estimation, as regularization will also be needed
 - In LogR we should minimize $-\ell(w_0, w; S) + \frac{\alpha}{2} \|w\|^2$

How to Handle Posterior Probabilities

- If possible, we don't want labels as model outputs but **posterior probabilities**
- Most models give them as pairs

$$(\hat{P}(0|x), \hat{P}(1|x)) = (\hat{P}(0|x), 1 - \hat{P}(0|x))$$

- In principle we would decide 1 if $\hat{P}(1|x) > 0.5$ and viceversa, but this may be too crude
- It may be advisable to set a **decision threshold** $0 < \theta < 1$ and decide 1 if $\hat{P}(1|x) > 1 - \theta$ and 0 if $\hat{P}(1|x) < \theta$
- For **imbalanced** problems where $\pi_0 \gg \pi_1$ (usually the interesting ones) we would have $\hat{P}(1|x) \simeq 0$ for most x
 - In this case we may choose another $\theta \ll 0.5$ and **suggest** 1 if $\hat{P}(1|x) > \theta$

Takeaways on Classification

- 1 We have introduced the classification problem as one of **computing posterior probabilities**
- 2 We have found the **optimal Bayes classifier** and **approximated it by k -NN**
- 3 We have introduced **Logistic Regression** and the numerical minimization of its (minus) log-likelihood
- 4 We have introduced several measures of **classifier performance**
- 5 We have reviewed some **practical issues** of classification

Rosenblatt's Perceptrons

1 Machine Learning Basics


- Basic Modeling
- Regression Basics
- Bias, Variance and Cross Validation
- Basic Classification
- Logistic Regression

2 Multilayer Perceptrons

- Rosenblatt's Perceptrons**
- Classical MLPs
- Neural Network Classification
- MLP Regularization
- Unconstrained Smooth Optimization
- Computational Costs of MLPs

Basic Neural Models


- Basic model: Ramón y Cajal's neuron (1900)



- Basic behavior: the neuron either fires or stays at rest depending basically on its inputs
- This translates into an electrical nerve impulse or **action potential**
- The brain has about 10^{11} neurons
 - Each one has about 7,000 connections
 - These connections are often **recurrent**

Hodgkin–Huxley


- They developed (circa 1935) the first electro-physiological model to describe the generation and propagation of action potentials in neurons
- They used the giant squid axon for this



From Wikipedia's [Squid Giant Synapse](#)

Electronic Neuron

- Electronic version: McCulloch–Pitts (1940)



- Since $I = \frac{V}{R} = wV$, with w the conductance, the McCulloch–Pitts neuron output is

$$H \left(\sum_{j=1}^d w_j V_j + I \right) = V$$

with the Heaviside function H ensuring a 0–1 output

- The **perceptron** is a mathematical formulation of the McC–P neuron where **the conductances are to be learned**

Rosenblatt's Perceptron

- Given a sample $S = \{(x^p, y^p)\}$ with $y^p = \pm 1$, **Rosenblatt's Perceptrons** (PCPs) are linear machines $w \cdot x$ such that

$$\begin{aligned} w \cdot x^p &> 0 & \text{if } y^p &= 1 \\ w \cdot x^p &< 0 & \text{if } y^p &= -1 \end{aligned}$$

- This can be seen as a **homogeneous** classification problem
- In more compact form we want for all p

$$y^p w \cdot x^p > 0$$

- If such a **separating** w exists, it can be computed in many ways
- But Rosenblatt's goal was to **learn** w , i.e., to arrive at a separating w by repeatedly examining the (x^p, y^p) and adjusting w if necessary

Rosenblatt's Delta Rule

- The **Delta Rule** algorithm

Start with $w = 0$

While stopping condition not met:

 Get a new pattern $(x^p, y^p) = (x^{p(t)}, y^{p(t)})$

 if $y^p w \cdot x^p \leq 0$ then:

$w = w + y^p x^p$

- Writing w for the new vector and w' for the previous one, notice that

$$y^p w \cdot x^p = y^p w' \cdot x^p + \|x^p\|^2 > y^p w' \cdot x^p$$

Novikov's Theorem

- It states that if the sample S is **linearly separable**, the Delta rule yields a separating w in a **finite number of steps**
 - I.e., it **learns**!!
- More precisely, if for some separating hyperplane w with $\|w\| = 1$ we have $\min_p y^p w \cdot x^p = \gamma > 0$, then **a separating w' can be found after at most $T = R^2/\gamma^2$ wrong classifications, where $R = \max_p \|x^p\|$**
- Notice that γ is the **margin** of the w hyperplane. Thus, if


$$\gamma^* = \max_{\|w\|=1} m(w, S) = \max_{\|w\|=1} \min_p y^p w \cdot x^p$$

is the **maximum margin**, we will need at most $R^2/(\gamma^*)^2$ iterations to train a Perceptron

- The maximum margin hyperplane gives the best (smallest) Novikov bound
- And problems with smaller margins require more iterations


Margins and Generalization

- Intuitively, when the margin is large, the problem is easier



Distance to a Line

- The margin of w is just the smallest distance between sample points and the hyperplane defined by w



- This extends to the multidimensional case: if $\|w\| = 1$ and $b = 0$, the distance of x to the hyperplane defined by w is $|w \cdot x| = yw \cdot x$
- Thus the Delta Rule **corrects weight vectors with negative margins**

Novikov's Proof Sketch

- Assume x^t is the t -th pattern such that $y^t w^{t-1} \cdot x^t \leq 0$
- Then $\|w^t\|^2 = \|y^t x^t + w^{t-1}\|^2 \leq \|x^t\|^2 + \|w^{t-1}\|^2$ and, thus,

$$\|w^t\|^2 \leq \|x^t\|^2 + \dots + \|x^1\|^2 + \|w^0\|^2 \leq tR^2$$

if $w^0 = 0$ and $\|x^p\| \leq R$ for all p

- Also, if w' is a 1-norm separating hyperplane i.e., $\min_p y^p w' \cdot x^p = \gamma > 0$ for all p , we have

$$w^t \cdot w' = (y^t x^t + w^{t-1}) \cdot w' \geq \gamma + w^{t-1} \cdot w' \geq \dots \geq t\gamma + w^0 \cdot w'$$

- Putting both together we have for all t for which x^t is erroneous


$$t\gamma \leq w^t \cdot w' \leq \|w^t\| \leq \sqrt{t}R$$

which implies $t \leq R^2/\gamma^2$

- Thus, the Delta Rule will stop after at most R^2/γ^2 iterations

Problems with Perceptrons

- Rosenblatt's Perceptron caused a big stir on the feasibility of **Conexionist Artificial Intelligence** but soon its first winter came
- Observation 1: A PCP can learn (obviously) only linearly separable problems; for instance, it can learn the AND and OR predicates




(a) AND



(b) OR

- But a **PCP cannot learn the XOR predicate**



Dichotomies

- Observation 2: for PCPs to be useful, **linearly separable problems must be frequent**
- Q: how frequent are they?
- A **dichotomy** for $S = \{x^1, x^2, \dots, x^N\}$ is any separation of S in two classes
- It corresponds to any assignment of the ± 1 values to the x^p
- The total number of dichotomies is 2^N
- Thus, the **total number of 2-class problems over S is 2^N**


Cover's Theorem

- We have to compare the total number of dichotomies with the number of dichotomies that are **linearly separable**
- The points in a sample S are said to be in **general position** if there are not $d + 1$ patterns in a $d - 1$ -dimensional hyperplane
- **Theorem** If S is in general position, the number $L(N, d)$ of linearly separable dichotomies is

$$L(N, d) = \left\{ \begin{array}{ll} 2^N & \text{si } N \leq d + 1 \\ 2 \sum_{i=0}^d \binom{N-1}{i} & \text{si } N \geq d + 1 \end{array} \right\}$$


Points in General Position

- Consider $d = 2$, $3 = d + 1$ points and a $1 = d - 1$ -dimensional hyperplane



Points Not in General Position

- Consider now $d = 2$ and $3 = d + 1$ points **not** on a $1 = d - 1$ -dimensional hyperplane (i.e., a line)



Thus ...

- Notice that for d fixed,

$$\frac{L(N, d)}{2^N} \rightarrow 0$$

as $N \rightarrow \infty$

- In practice we can expect $N \gg d$ and the fraction of separable dichotomies will be very small
- Thus, linear PCPs will not be useful
- But a **one hidden layer** PCP can solve XOR and, in fact, **separate any convex region** from the rest of the space
- And a **two hidden layer** PCP can separate any **polyhedral region** and hence **solve any classification problem**
- The same is essentially true for regression problems
- But: **no algorithm is known to learn these 2-hidden layer PCPs**

Why Not Try to Minimize Some Loss?

- Rosenblatt's Perceptrons are natural **online** machines:
 - Patterns appear sequentially and each may change the current weights
- We can also define a perceptron loss over a given sample as

$$L(w|S) = - \sum_{\{p: y^p w \cdot x^p < 0\}} y^p w \cdot x^p$$

- The local error would thus be $-y^p w \cdot x^p$ when $y^p w \cdot x^p < 0$ and 0 otherwise
- Since $\nabla_w(-y^p w \cdot x^p) = -y^p x^p$, the Delta Rule can be seen as a kind of local gradient descent on the loss $L(w|S)$:

$$w^t = w^{t-1} + y^{p(t)} x^{p(t)} = w^{t-1} - \nabla_w \left(-y^{p(t)} x^{p(t)} \cdot w \right)$$

- However $L(w|S)$ is not differentiable and it is not clear what to make of $\nabla_w L(w|S) \dots$

What Can We Do?

- First option: work with linear models but somehow ensure that $d \gg N$
 - Transform the original d -dimensional features x in new D -dimensional ones such that $D \gg d$
 - We overcome then that, in practice, we usually have $N \gg d$
 - This is the goal of **kernel-based Support Vector Machines**
- Second option: build trainable PCPs replacing the Heaviside function
 - Work with differentiable PCP transfer function
 - Transform PCP training into a differentiable optimization problem
 - This is what we will do next
- Before doing so, note that **Rosenblatt's PCPs can be seen as the first Machine Learning model**
- And modern ML started around 1985 along the PCP's ideas

Classical MLPs

1 Machine Learning Basics

- Basic Modeling

- Regression Basics

- Bias, Variance and Cross Validation

- Basic Classification

- Logistic Regression

2 Multilayer Perceptrons

- Rosenblatt's Perceptrons

- Classical MLPs**

- Neural Network Classification


- MLP Regularization

- Unconstrained Smooth Optimization

- Computational Costs of MLPs

MLP Architecture

- General layout:
 - An input layer (input)
 - One or several hidden layers
 - One output layer
- Feedforward connections only



- Example: [TensorFlow Playground](#)

MLP Connections

- **No feedback or lateral connections**
- **Fully connected layers**
- **Linear weight connections** between units and (usually) **non linear activations** inside each unit
 - Combined effect of successive layers: potentially highly non-linear transformation
- General processing: layered and feedforward
- In practice (1990s), one hidden layer and only sometimes two
- Later (around 2010): Deep Networks with “many” (from 3 to 10) layers
- Now: almost anything!

Unit Activation and Output

- The **activations** a_i^h of a unit in layer h receives the **outputs** o_j^{h-1} from processing in the previous layer

$$a_i^h = \sum_{j=1}^{n_{h-1}} w_{ij}^h o_j^{h-1} + b_i^h,$$

- In matrix/vector form:

$$\mathbf{a}^h = \mathbf{W}^h \mathbf{o}^{h-1} + \mathbf{b}^h$$

- Output** o_i^h of a unit: **non linear processing** $o_i^h = \varphi(a_i^h)$ of its activation
- In vector form:

$$\mathbf{o}^h = \varphi(\mathbf{a}^h),$$

where φ is applied over each unit

Activation Functions

- Choices for f :
 - Heaviside (in the very first Rosenblatt's Perceptrons): $\varphi(a) = 0$ if $a \leq 0$, $\varphi(a) = 1$ if $a > 0$
 - **Identity/linear**: $\varphi(a) = a$
 - **Sigmoid**:

$$\varphi(a) = \sigma(a) = \frac{1}{1 + e^{-a}}$$


- **Hyperbolic tangent**:

$$\varphi(a) = \tanh(a) = \frac{e^a - e^{-a}}{e^a + e^{-a}}$$

- **Rectified Linear Units (ReLU)**: $\varphi(a) = r(a) = \max(0, a)$

Sigmoid, Hyperbolic Tangent, ReLUs

- Sigmoid and tanh: smooth version of Heaviside step function
- For ReLUs $r'(x)$ is either 0 or 1 (hoping $x = 0$ never happens!!)
 - Many gradient elements will go to 0 and many units will have constant activations no matter their input



From [Stanford's UFLDL Tutorial](#)

- Usual activation choices:
 - For hidden units: hyperbolic tangent and, currently, ReLU
 - Linear outputs for regression, sigmoid/soft max outputs for classification

The Forward Pass

- We need first to compute an MLP's outputs through the **forward pass**
- At each layer h we compute two variable sets
 - The linear **activations** $a^h = W^h o^{h-1} + b^h$ with o^{h-1} the previous layer's output and W^h, b^h the **weights** and **biases** connecting layer $h - 1$ with layer h
 - The nonlinear **outputs** $o^h = \varphi(a^h)$,
- In unit-wise terms: $a_i^h = \sum w_{ij}^h o_j^{h-1} + b_i^h$, $o_i^h = \varphi(a_i^h)$
- Concatenating these operations we arrive at the outputs $\hat{y} = f(x; \mathcal{W})$ of the network, with $\mathcal{W} = (W^1, b^1, \dots, W^H, b^H)$ the network's weight set

The Simplest MLP I

- The Single Hidden Layer (SHL) MLP
 - D inputs (determined by the problem at hand)
 - One hidden layer with H units (number to be chosen) and \tanh (or what you prefer) activation
 - One linear or sigmoid output (according to the problem at hand)
- Input–hidden processing: denoting inputs by x and the hidden unit output as o^H ,

$$o_h^H = \tanh \left(b_h^H + \sum_{j=1}^D w_{hj}^H x_j \right)$$

- In matrix/vector form: $o^H = \tanh (W^H x + b^H)$

The Simplest MLP II

- Hidden–output processing: since targets are 1–dimensional, we have for the outputs \hat{y}

$$\hat{y} = \sum_{h=0}^H w_h^O o_h^H + b^O,$$

- In vector form: $\hat{y} = w^O \cdot o^H + b^O$
- That is, a **linear model** on the last hidden layer outputs o^H
- Global process:

$$\hat{y} = f(x; w^O, W^H, b^O, b^H) = b^O + \sum_h w_h^O \tanh \left(b_h^H + \sum_j W_{hj}^H x_j \right)$$

- Or in matrix/vector form

$$\hat{y} = f(x; w^O, W^H, b^O, b^H) = b^O + w^O \cdot \tanh \left(b^H + W^H x \right)$$

- Or in Python `y_pred = np.tanh(x.dot(w_H.T) + b_H).dot(w_O.T) + b_O`

MLP Error Function

- MSE is the standard error function for regression MLPs

$$\begin{aligned}e(\mathcal{W}) &= \frac{1}{2}E_{x,y} [(y - f(x; \mathcal{W}))^2] = E_{x,y} [e^\ell(x, y; \mathcal{W})] \\&= \int e^\ell(x, y; \mathcal{W}) p(x, y) dx dy\end{aligned}$$

with $e^\ell(x, y; \mathcal{W})$ denotes the **local error**

$$e^\ell(x, y; \mathcal{W}) = \frac{1}{2}(y - \hat{y})^2 = \frac{1}{2}(y - f(x; \mathcal{W}))^2$$

Minimizing the MSE

- The general idea would be to obtain the optimal \mathcal{W}^* as a solution of $\nabla e(\mathcal{W}) = 0$, where

$$\begin{aligned}\nabla e(\mathcal{W}) &= E_{x,y} [\nabla_{\mathcal{W}} e^{\ell}(x, y; \mathcal{W})] \\ &= E_{x,y} [\nabla_{\mathcal{W}} f(x; \mathcal{W})(f(x; \mathcal{W}) - y)]\end{aligned}$$

for we have

$$\begin{aligned}\nabla_{\mathcal{W}} e^{\ell}(x, y; \mathcal{W}) &= -(y - f(x; \mathcal{W})) \nabla_{\mathcal{W}} f(x; \mathcal{W}) \\ &= \nabla_{\mathcal{W}} f(x; \mathcal{W})(f(x; \mathcal{W}) - y)\end{aligned}$$

- This suggests to get $\nabla_{\mathcal{W}} f(x; \mathcal{W})$ but it is much better to focus on ∇e^{ℓ} , and
 - We will compute first ∇e^{ℓ} (and then $\nabla e = E[\nabla e^{\ell}]$)
 - We will exploit it to build MLPs through **optimization methods**

The Forward Pass (again)

- Recall that we compute an MLP's outputs through the **forward pass**
- At each layer h we compute (and **store**) first the linear **activations** $a^h = w^h o^{h-1} + b^h$ with o^{h-1} the previous layer's output and W^h, b^h the **weights** and **biases** connecting layer $h - 1$ with layer h
- We then compute (and store) the nonlinear **outputs** $o^h = \varphi(a^h)$,
- In unit-wise terms:

$$a_i^h = \sum w_{ij}^h o_j^{h-1} + b_i^h, \quad o_i^h = \varphi(a_i^h)$$

- And we also have

$$\frac{\partial a_i^h}{\partial w_{ij}^h} = o_j^{h-1}, \quad \frac{\partial o_i^h}{\partial a_i^h} = \varphi'(a_i^h)$$

Computing the Gradient

- The key tool is the **chain rule** that we will apply backwards from the output layer
- If w_{ij}^h is the general weight connecting unit j in layer $h - 1$ to unit i in layer h , we have

$$\frac{\partial e^\ell}{\partial w_{ij}^h} = \frac{\partial e^\ell}{\partial a_i^h} \frac{\partial a_i^h}{\partial w_{ij}^h} = \frac{\partial e^\ell}{\partial a_i^h} o_j^{h-1} = \delta_i^h o_j^{h-1}$$

- That is, we have an easy to compute term $\frac{\partial a_i^h}{\partial w_{ij}^h}$, and a more complicated but crucial one $\frac{\partial e^\ell}{\partial a_i^h}$
- We refer to $\frac{\partial e^\ell}{\partial a_i^h} = \delta_i^h$ as the **generalized error** at unit i of layer h

Gradient at the Output Layer

- Let's assume that we have layers $1, 2, \dots, H-1, H$, with H the output layer
- For simplicity, consider a regression problem with linear outputs
- In the output H layer we have

$$e^\ell = \frac{1}{2}(y - \hat{y})^2 \quad \text{and} \quad \hat{y} = o^H = a^H = \sum_j w_j^H o^{H-1}$$

- Thus, the partial derivative is straightforward

$$\delta^H = \frac{\partial e^\ell}{\partial a^H} = \frac{\partial e^\ell}{\partial \hat{y}} = \hat{y} - y$$

i.e., the generalized error δ^H is here the **standard error**

- Therefore, we have for the output layer H

$$\frac{\partial e^\ell}{\partial w_j^H} = (\hat{y} - y) \frac{\partial a^H}{\partial w_j^H} = (\hat{y} - y) o_j^{H-1}$$

- Now it is (quite) easy to extend this to other layers

Backprop for Other Layers

- Let w_{ij}^h be the weight connecting unit j in layer $h - 1$ with unit i in layer h ; recall that then

$$\frac{\partial e^\ell}{\partial w_{ij}^h} = \frac{\partial e^\ell}{\partial a_i^h} \frac{\partial a_i^h}{\partial w_{ij}^h} = \frac{\partial e^\ell}{\partial a_i^h} o_j^{h-1}$$

- And if we have already computed the generalized errors $\delta_k^{h+1} = \frac{\partial e^\ell}{\partial a_k^{h+1}}$ for layer $h + 1$, we have


$$\begin{aligned} \frac{\partial e^\ell}{\partial a_i^h} &= \sum_k \frac{\partial e^\ell}{\partial a_k^{h+1}} \frac{\partial a_k^{h+1}}{\partial a_i^h} = \sum_k \delta_k^{h+1} \frac{\partial a_k^{h+1}}{\partial a_i^h} \\ &= \sum_k \delta_k^{h+1} \frac{\partial a_k^{h+1}}{\partial o_i^h} \frac{\partial o_i^h}{\partial a_i^h} = \left[\sum_k \delta_k^{h+1} w_{ki}^{h+1} \right] \varphi'(a_i^h) \end{aligned}$$

Backprop and Deep Networks

- The previous formulae work for any number of hidden layers
 - We can work with **deep** MLPs
- Moreover, the chain rule can be automated and modern NN packages compute gradients automatically for very general NN architectures
 - Basic tool: automatic symbolic differentiation on computational graphs
 - This is one of the key reasons of the great success of deep neural nets
- Some examples: convolutional nets, residual nets, recurrent NNs, ...

Automating Backprop

- To deal with such general networks requires **automating backprop**
- This is done upon the network's **computational graph**



From Zhang et al., **Dive into Deep Learning**

- It corresponds to a SHL with a loss $J = L + s$
- The circles contain differentiable operators
- The shadowed squares contain the variables against which differentiation takes place

Neural Network Classification

1 Machine Learning Basics

Basic Modeling

Regression Basics

Bias, Variance and Cross Validation

Basic Classification

Logistic Regression

2 Multilayer Perceptrons

Rosenblatt's Perceptrons

Classical MLPs

Neural Network Classification

MLP Regularization

Unconstrained Smooth Optimization

Computational Costs of MLPs

Logistic Regression (again)

- Recall that our Logistic Regression model gave us estimate of $P(0|x; w_0, w)$, $P(1|x; w_0, w)$
- We then wrote the likelihood of a given sample as

$$P(Y|X; w_0, w) = \prod_{p=1}^N P(y^p|x^p; w_0, w)$$

- And then derived the optimum parameters w_0^* , w^* as

$$w_0^*, w^* = \arg \max_{w_0, w} P(Y|X; w_0, w)$$

- Notice that we can see the output of a logistic regression model as that of an **NN with no hidden layers and sigmoid outputs**

Neural 2-class Classification

- We can follow the same path for a 2 class classification problem changing that output for the one

$$f(x; \mathcal{W}) = w_0 + w \cdot \Phi(x, \mathcal{W}^H)$$

with w_0, w the weights from the last hidden layer (LHL) to the output and $\Phi(x, \mathcal{W}^H)$ the map from inputs to the LHL outputs

- And now derive the optimum parameters \mathcal{W}^* as

$$\begin{aligned}\mathcal{W}^* &= \arg \min_{\mathcal{W}} - \sum_{p=1}^N \log P(y^p | x^p; \mathcal{W}) \\ &= \arg \min_{w_0, w, \mathcal{W}^H} - \sum_p y^p (w_0 + w \cdot \Phi(x^p, \mathcal{W}^H)) + \\ &\quad \sum_p \log(1 + e^{w_0 + w \cdot \Phi(x^p, \mathcal{W}^H)})\end{aligned}$$

- Forward and backpropagation are also straightforward here
- We change things a little bit for K class problems

MLPs for K -class Classification

- We consider an input layer and a number of hidden layers
- Targets are the 1-hot encodings of class labels, so we use K outputs
- We want the MLP's k -th output to estimate the posterior $P(k|x)$
- The natural output layer activation is thus the **softmax function**

$$\sigma_j(o) = F_j(x; \mathcal{W}) = \frac{e^{w_0^j + w^j \cdot o}}{\sum_1^K e^{w_0^k + w^k \cdot o}}$$

- Here o are the last hidden layer outputs and w^j , w_0^j are the weight and bias from the last hidden layer to the network's j output
- Moreover $\sum_j F_j(x; \mathcal{W}) = 1$ and we assume $P(j|x) \simeq F_j(x; \mathcal{W})$
- For two classes this becomes

$$\sigma_1(o) = \frac{e^{w_0^1 + w^1 \cdot o}}{e^{w_0^0 + w^0 \cdot o} + e^{w_0^1 + w^1 \cdot o}} = \frac{1}{1 + e^{w_0^0 - w_0^1 + (w^0 - w^1) \cdot o}} = \frac{1}{1 + e^{\omega_0 + \omega \cdot o}}$$

- We thus get the **sigmoid activation** of Logistic Regression

Multiclass Log–Loss NN Classification

- We change slightly the preceding for a K class classification problem
- We define an MLP $F(x; \mathcal{W}) \in \mathbf{R}^K$ with softmax outputs
 - Recall that $0 \leq F_k(x; \mathcal{W}) \leq 1$ and $\sum_k F_k(x; \mathcal{W}) = 1$
 - And also we assume that $P(k|x) \simeq F_k(x; \mathcal{W})$
- We label now the multiclass targets via the 1–hot encoding, with the target of class k being the vector $\mathbf{e}_k = (0, \dots, \underbrace{1}_k, \dots, 0)$
- Then, if $x \in C_k$, its label is $y = (y_1, \dots, y_K)^t$ is \mathbf{e}_k , with $y_k = 1$, and

$$P(k|x) = P(k|x)^1 = P(k|x)^{y_k} = \prod_{c=1}^K P(c|x)^{y_c} \simeq \prod_{c=1}^K F_c(x; \mathcal{W})^{y_c}$$

The Sample Likelihood

- We have a sample $S = \{x^p, y^p\}$ with (x^p, y^p) in class $c(p)$ and, hence, $y^p = e_{c(p)}$
- We also consider a posterior model $P(c|x, w) = F_c(x, \mathcal{W})$, the probability of getting $S = (X, Y)$ is

$$\begin{aligned} P(Y|X; \mathcal{W}) &= \prod_1^N P(c(p)|x^p; \mathcal{W}) = \prod_1^N \prod_{c=1}^K P(c|x^p; \mathcal{W})^{y_c^p} \\ &\simeq \prod_1^N \prod_{c=1}^K F_c(x; \mathcal{W})^{y_c^p} \end{aligned}$$

- Thus, we estimate \mathcal{W} by maximizing the sample's **likelihood**

$$\mathcal{P}(Y|X; \mathcal{W}) = \prod_{p=1}^N \prod_{c=1}^K F_c(x^p; \mathcal{W})^{y_c^p}$$

The Cross Entropy Loss

- As before, we will work with the sample's log-likelihood, i.e.

$$\begin{aligned}\ell(\mathcal{W}; S) &= \log \mathcal{P}(Y|X; \mathcal{W}) = \sum_{p=1}^N \sum_{c=1}^K \log \left(F_c(x^p; \mathcal{W})^{y_c^p} \right) \\ &= \sum_{p=1}^N \sum_{c=1}^K y_c^p \log F_c(x^p; \mathcal{W})\end{aligned}$$

- The **cross-entropy** loss is now simply the negative log-likelihood $-\ell(\mathcal{W}; S)$, i.e.,

$$L(\mathcal{W}) = -\ell(\mathcal{W}; S) = - \sum_{p=1}^N \sum_{c=1}^K y_c^p \log F_c(x^p; \mathcal{W})$$

- It is now straightforward to carry this into a NN setting

Takeaways on Classical MLPs

- ① They have a layered structure with outputs computed in a **forward pass** using differentiable activations
- ② Usual activations: ReLUs, sigmoid, tanh, linear
- ③ MSE is the usual regression cost; we use a variant of cross entropy in classification
- ④ The error function gradients are computed by **backpropagation** of generalized errors
- ⑤ Backprop is basically a very simple procedure than can be **largely automated**
- ⑥ Regression and classification problems fit very well in this setting

MLP Regularization

1 Machine Learning Basics

- Basic Modeling

- Regression Basics

- Bias, Variance and Cross Validation

- Basic Classification

- Logistic Regression

2 Multilayer Perceptrons

- Rosenblatt's Perceptrons

- Classical MLPs

- Neural Network Classification

- MLP Regularization**

- Unconstrained Smooth Optimization

- Computational Costs of MLPs

MLPs and Universal Approximation

- In regression we assume $y = \phi(x) + n$ and our goal is to get $f \simeq \phi$

- We say that $\mathcal{F} = \{f(x; \mathcal{W})\}$ is a **Universal Approximation Family** over a domain \mathcal{R} if


For any $\epsilon > 0$ and any reasonable ϕ , we can find an $f(x; \mathcal{W}_{\phi, \epsilon})$ s.t.

$$e(\mathcal{W}_{\phi, \epsilon}) = \int (\phi(x) - f(x; \mathcal{W}_{\phi, \epsilon}))^2 p(x) dx \leq \epsilon$$

- Notice that Universal Approximation is just what we need in regression
- In fact a **Single Hidden Layer MLP with enough hidden units is an effective universal approximator**
- But it also implies the risk of **overfit** whose control is crucial

Overfitting in MLPs

- Since MLPs are a UAF, they can also approximate the noise in the sample
 - Given $S = \{(x^p, y^p)\}$ if we allow enough hidden units in a SHL MLP we can arrive to a w^* s.t. $y^p = f(x^p; w^*)$
 - We get thus a sample error $\hat{e}(w^*) = 0$ but possibly with a very high generalization error
- I.e., MLPs may have very small bias but possibly large variance



(Ph.D. Thesis of Carlos Alaíz)

Regularization vs Overfitting

- Why is there overfitting?
 - Because we may end up having too many weights with respect to sample size
 - Because we allow these weights to explore the entire weight space
- We can avoid this wandering if we limit w 's growth, for which we add a **regularization** term $g(\|w\|)$ to $e(w)$ with $g(r)$ increasing
- Working with $e_R(w) = e(w) + g(\|w\|)$ we have to **balance** the minimization of $e(w)$ and that of $g(\|w\|)$
- This balanced learning results in better generalization

L_2 Regularization

- The simplest regularization procedure adds a quadratic penalty to the square error e

$$e_R(w) = e(w) + \frac{\lambda}{2} \|w\|^2,$$

with λ the **weight decay** factor

- Also known as Tikhonov's regularization (or Ridge Regression for linear models)
- The desired effect is to constrain the evolution of w :
 - In fact, the regularized loss is essentially the **Lagrangian of the constrained problem**

$$\min_w e(w) \text{ subject to } \|w\|^2 \leq \rho, \quad \rho > 0$$

- The gradient becomes $\nabla e_R(w) = \nabla e(w) + \lambda w$ and gradient descent becomes

$$w^{k+1} = w^k - \rho_k (\nabla_w e(w^k) + \lambda w^k)$$

How to choose λ

- Again, the correct choice of λ is crucial
- A small $\lambda \ll 1$ results in a small regularization effect and overfitting risk appears
- A large $\lambda \gg 1$ causes learning to forget about $e(w)$ and the model will be essentially constant and will underfit
- Usually λ is chosen by using CV by exploring a discrete set of values λ_j ,
- The same can essentially be done for any other hyper-parameter:
 - Number of hidden layers and of hidden units?
 - Learning rate??
 - Minibatch size???
- But this may become unmanageable if we want to work with many (more than three?) hyperparameters

MLP Ensembles

- Recall that for MLPs $e(w)$ does not have a single minimum
- Moreover, the final MLP depends on the random initial w^0
- And mini-batch training adds extra randomness to the final model
- Therefore, **you may never get the same MLP twice!!!**
- But we can turn this to our advantage by
 - Starting from K independent initial weights and get K optimal weight sets w_k^*
 - Giving as the output the average $f_e(x) = \frac{1}{K} \sum_1^K f(x; w_k^*)$
- We expect outputs of the form $\hat{y}_k^p = y^p + \epsilon_k^p$ with the errors ϵ_k^p being independent
- Hence $\frac{1}{K} \sum_k \epsilon_k^p \simeq 0$ and $\frac{1}{K} \sum_k \hat{y}_k^p \simeq y^p$

Takeaways on MLP Regularization

- MLPs have a **high risk of overfitting**
- Thus, they must be **regularized** to avoid overfitting
 - The regularization hyperparameter is chosen through **cross validation**
- MLP training has two **random components**: the initial point and minibatch selection
 - Thus, each training will converge to a **different optimum**
- MLP **ensembles** can take advantage of this
 - They imply extra costs but ensembles are **embarrassingly parallelizable**

Unconstrained Smooth Optimization

1 Machine Learning Basics

Basic Modeling

Regression Basics

Bias, Variance and Cross Validation

Basic Classification

Logistic Regression

2 Multilayer Perceptrons

Rosenblatt's Perceptrons

Classical MLPs

Neural Network Classification

MLP Regularization

Unconstrained Smooth Optimization


Computational Costs of MLPs

Back to Optimization

- To build a ML model \equiv to **minimize a loss function**
- General optimization theory is a key tool in Machine Learning (ML)
- There are two optimization set ups in ML
 - **Unconstrained** optimization, slightly simpler and the one used for MLPs
 - **Constrained** optimization, wider and more complex
- In ML we have also to consider the optimization of differentiable and also non differentiable error functions
- MLP optimization: **unconstrained and differentiable**
- And also **batch**, i.e., over the entire sample, **mini-batch** over subsamples or **on line**, pattern by pattern

Gradient Descent

- We “stretch” the weight set \mathcal{W} into a vector w
- We recall that $-\nabla e(w)$ is the maximum descent direction
- First idea: to build a (hopefully convergent) sequence w^k iterating (small) steps along $-\nabla e(w^k)$



Gradient Descent II

- In more detail, we start from a random w^0 and compute

$$w^{k+1} = w^k - \rho_k \nabla_w e(w^k)$$

- ρ_k is the **learning rate** (LR)
- With a small ρ_k we ensure $e(w^{k+1}) < e(w^k)$ (although with possibly a very small descent)
- We can get a better iteration $w^{k+1} = w^k - \rho_k^* \nabla_w e(w^k)$ using a ρ_k^* given by

$$\rho_k^* = \arg \min_{\rho} e(w^k - \rho \nabla_w e(w^k));$$

this is known as **line minimization**

- These GD methods are called **first order methods** in part because they only use ∇e

One Dimensional Newton's Method

- Assume a quadratic function $q(w) = aw^2 + bw + c$, with $a > 0$, and a minimum at w^*
- We can reach w^* from any w with a step Δw such that

$$0 = q'(w + \Delta w) = 2a(w + \Delta w) + b$$

- We have thus $\Delta w = \frac{-b-2aw}{2a}$, that is,

$$w^* = w - \frac{2aw + b}{2a} = w - \frac{1}{q''(w)} q'(w)$$

- This leads to **Newton's method**: minimize a general f iteratively using steps

$$w^{k+1} = w^k - \rho_k \frac{1}{f''(w^k)} f'(w^k)$$

with ρ_k a suitable learning rate

One Dimensional Newton's Method II

- Notice how the $\frac{1}{f''(w_k)}$ acts as a **self adjusting learning rate**
- When the parabola is very sharp (i.e., $a \gg 1$), we will be close to the minimum and the descent can overstep it
- But then $f''(w_k)$ will be big and, hence, $\frac{1}{f''(w_k)}$ **moderates** the gradient step
- Conversely, when the parabola is very wide (i.e., $a \simeq f''(w_k) \ll 1$), the minimum will be far away and the gradient step small descent will
- But then $f''(w_k)$ will be rather small and, hence, $\frac{1}{f''(w_k)}$ **accelerates** the gradient step
- We can extend this to several dimensions

Multidimensional Newton's Method

- For a d dimensional w , the Taylor expansion of e at an optimum w^* is

$$e(w) \approx e(w^*) + \frac{1}{2}(w - w^*)^t \cdot \mathcal{H}(w^*) \cdot (w - w^*)$$

- $\mathcal{H}(w^*)$ is the **Hessian** of e at w^* and $\nabla e(w^*) = 0$
- It follows that $\nabla e(w) \approx \mathcal{H}(w^*) \cdot (w - w^*)$ and, therefore,

$$w^* \approx w - \mathcal{H}(w^*)^{-1} \nabla_w e(w)$$

- This suggest to derive the w^k by

$$w^{k+1} = w^k - \rho_k \mathcal{H}(w^k)^{-1} \nabla_w e(w^k)$$

which is known as **Newton's Method** (NM)

Variants of Newton's Method

- Theoretically NM converges very fast near w^* , but
 - Far from w^* convergence is not guaranteed
 - Moreover $\mathcal{H}(w^k)$ may not be invertible
 - Besides, computing $\mathcal{H}(w^k)$ is cumbersome and costly
- The **Gauss–Newton (GN)** approximation

$$\mathcal{H}(w) \simeq E[\nabla e(w) \nabla e(w)^T],$$

holds for any quadratic cost and simplifies the third problem

- The **Levenberg–Marquardt (LM)** method deals with the other two problems combining
 - Gradient descent “away” from w^*
 - Gauss–Newton “near” w^*

although “away” and “near” have to be properly addressed

One-dimensional GN Approximation

- Assume $e(w) = \frac{1}{2} \int (f(x; w) - y)^2 p(x, y) dx dy$; then

$$e'(w) = \int (f(x; w) - y) \frac{\partial f}{\partial w}(x, y) p(x, y) dx dy;$$

$$\begin{aligned} e''(w) = & \int \left(\frac{\partial f}{\partial w} \right)^2 p(x, y) dx dy \\ & + \int (f(x; w) - y) \frac{\partial^2 f}{\partial w^2}(x, y) p(x, y) dx dy \end{aligned}$$

- Near a minimum w^* we may expect $f(x; w) \simeq y$ and, therefore,

$$e''(w) \simeq \int \left(\frac{\partial f}{\partial w} \right)^2 p(x, y) dx dy > 0$$

- Thus, for square errors, we can use first derivatives to approximate $e''(w)$

General GN Approximation I

- In the general case we have

$$\nabla e(w) = E [\nabla f(x; w) (f(x; w) - y)]$$

- And, therefore,

$$\begin{aligned} \nabla^2 e(w) &= E [\nabla^2 f(x; w) (f(x; w) - y)] + \\ &\quad E [\nabla f(x; w) \nabla f(x; w)^T] \end{aligned}$$

- The second term is easy to compute once we have ∇f
- If $w \approx w^*$, $f(x; w) \approx y$; therefore $f(x; w) - y \approx 0$,
 - We can ignore the first, more complex, term

General GN Approximation II

- We arrive at $\nabla^2 e(w) \simeq E [\nabla f(x; w) \nabla f(x; w)^\tau]$ or, equivalently,

$$\mathcal{H}_{(i,j)(p,q)}(w) = \left(\frac{\partial^2 e}{\partial w_{ij} \partial w_{pq}}(w) \right) \simeq \left(E \left[\frac{\partial f}{\partial w_{pq}} \frac{\partial f}{\partial w_{ij}} \right] \right)_{(i,j)(p,q)}$$

- $\mathcal{J} = E [\nabla f(x; w) \nabla f(x; w)^\tau]$ is **Fisher's information matrix**
 - \mathcal{J} is **semidefinite positive**
 - And thus not necessarily invertible
- Often only its diagonal is considered and we have


$$\mathcal{H}_{(i,j)(i,j)}(w) \simeq E \left[\left(\frac{\partial f}{\partial w_{ij}} \right)^2 \right]$$

Advanced Optimization

- There are many more proposals in unconstrained optimization
- The **Conjugate Gradient** (CG) and **Quasi-Newton** (QN) methods are important in classical MLP training
- The basic idea in CG is to replace gradient descent directions $g_k = -\nabla e(w_k)$ with new conjugate directions that try to keep somehow the previous “good directions”
- The basic idea in QN is to iterate as in NM but with simple approximations \mathcal{A}_k to $\mathcal{H}^{-1}(w^k)$ that converge to $\mathcal{H}^{-1}(w^*)$
- When training “small” NNs the **Limited-memory Broyden-Fletcher-Goldfarb-Shanno** (L-BFGS) QN variant is often used

Accelerating Gradient Descent

- A possibility on the error surface of a NN is to have many long, narrow ravines
 - Gradients bounce in the narrow section, but may be very small along the wider one



Momentum

- **Momentum** pushes them forward in the wider section (but we have to cope with a new parameter)
- Momentum tries to maintain descent's inertia with a term $\Delta^k = w^k - w^{k-1}$, i.e.,

$$w^{k+1} = w^k - \rho_k \nabla_w e(w^k) + \mu_k \Delta^k$$

- More or less, what a **heavy ball** would do
- The goal is to keep w^k advancing in “plateaux”, i.e., small gradient zones
- Momentum can be seen as a crude approximation of a CG step
- Nice explanation at [Why Momentum Really Works](#)

Nesterov's Accelerated Gradient

- Let's rewrite momentum in two steps
 - 1 Define $\Delta^{k+1} = -\rho_k \nabla_w e(w^k) + \mu_k \Delta^k$ and
 - 2 Apply

$$w^{k+1} = w^k + \Delta^{k+1} = w^k - \rho_k \nabla_w e(w^k) + \mu_k \Delta^k$$

- Nesterov's Accelerated Gradient** is a variant of this

$$\begin{aligned}\tilde{\Delta}^{k+1} &= -\rho_k \nabla_w e(w^k + \mu_k \tilde{\Delta}^k) + \mu_k \tilde{\Delta}^k; \\ w^{k+1} &= w^k + \tilde{\Delta}^{k+1}\end{aligned}$$

- In convex optimization it improves GD and is often highly effective in Deep Network training

Adam

- Adam is currently the most widely used gradient-descent method for deep NN training
- At each step t Adam uses a new random mini-batch to
 - Update **exponentially smoothed** averages m_t of the gradient g_t and v_t of the the squared gradient $g_t^2 = g_t \odot g_t$ as

$$m_t = \beta_1 m_{t-1} + (1 - \beta_1) g_t, \quad v_t = \beta_2 v_{t-1} + (1 - \beta_2) g_t^2;$$

- Since we have

$$E[m_t] \simeq (1 - \beta_1^t) E[g_t], \quad E[v_t] \simeq (1 - \beta_2^t) E[g_t^2],$$

we compute bias corrections \hat{m}_t, \hat{v}_t as

$$\hat{m}_t = \frac{1}{1 - \beta_1^t} m_t, \quad \hat{v}_t = \frac{1}{1 - \beta_2^t} v_t;$$

- Update weights as $W_t = W_{t-1} - \alpha \frac{\hat{m}_t}{\sqrt{\hat{v}_t + \epsilon}}$

Understanding Adam

- Default values $\alpha = 0.001$, $\beta_1 = 0.9$, $\beta_2 = 0.999$, and $\epsilon = 10^{-8}$ usually work fine
- What is going on in Adam?
- At each step we work with estimates of the average local gradient:

$$\hat{m}_t \simeq E[\nabla_w f]; \quad \hat{v}_t \simeq \sigma(g_t) \simeq E[(\nabla_w f)^2]$$

- Thus, we can first see the Adam iterations

$$W_t = W_{t-1} - \alpha \frac{\hat{m}_t}{\sqrt{\hat{v}_t} + \epsilon}$$

as a kind of “normalized” gradient descent

- In more detail, we expect

$$g_t^2 = E[(\nabla_w f)^2] \simeq E[\nabla_{w^2}^2 f] = \mathcal{H},$$

we can see Adam as a “dampened” variant of diagonal Gauss–Newton’s steps

When to Stop Training

- Typically the $e(\mathcal{W}^k)$ error diminishes towards an asymptotic minimum
 - If many units are used, we may arrive to 0 error, which usually implies **overfitting**
- First solution: to use a separate **validation subset** V and stop training when the error in V , i.e., the **validation error** starts growing
 - This **early stopping** is available in the main NN packages
 - But what do we do for small samples?
- Second solution: **get a good regularization**
 - Now training stops because of reasons such as computational cost, but not because of overfitting risk
- We must combine both, to save on training time and because overfitting is a serious MLP risk

Takeaways on MLP Optimization

- **Gradient descent** is the simplest method but possibly also the slowest
- **Momentum** can be used to speed it up
- **Newton's method** is much faster but may be very costly and difficult to apply in full form
- All the previous methods require the (usually tricky) selection of a **learning rate**
- Second order methods such as Conjugate Gradient and **Quasi-Newton** avoid learning rates and are more efficient but costlier
- **Limited Broyden-Fletcher-Goldfarb-Shanno** (L-BFGS) is currently the option of choice for "small" MLPs

Big Data

1 Machine Learning Basics

- Basic Modeling

- Regression Basics

- Bias, Variance and Cross Validation

- Basic Classification

- Logistic Regression

2 Multilayer Perceptrons

- Rosenblatt's Perceptrons

- Classical MLPs

- Neural Network Classification

- MLP Regularization

- Unconstrained Smooth Optimization

- Computational Costs of MLPs**

MLPs and Big Data?

- Many Vs in Big Data: **Volume**, **Velocity**, Variety, Veracity, Value, ...
- Velocity: information flows in data streams that require fast processing and feed back
 - MLPs are rather fast
 - Less than linear models but much more so than SVMs
- Volume is probably the greatest attractive of Big Data
 - Huge samples and/or very large pattern dimension
 - Large impact in model training

Volume in MLPs

- Parallelism is the first answer to Big Volume
- On a standalone machine it may be
 - Passive: let the SO distribute work among several CPU cores, or use GPUs with low level parallelized libraries such as Linpack or BLAS for linear algebra
 - Active: explicitly exploit a problem's parallelism programming an algorithm in, say, OpenMP
- It is easy to passively parallelize the training of several MLPs
 - But memory costs may multiply
- Training an isolated MLP is not CPU parallelizable:
 - There is a sequential layer dependence in Backprop
 - Each training iteration depends on the previous one
 - But matrix-vector multiplications can exploit GPU parallelism

Handling Huge Sample Training

- Two consequences of a large N are
 - Sample doesn't fit in memory and we have to split it somehow
 - Training gets "lost", for in the global gradient $\nabla e = E[\nabla e^\ell]$ we average many local gradients that may cancel each other out
- The first problem has been always present in fields such as analog signal filtering
- Solution for both: adaptive or **on line**, i.e., pattern by pattern, weight updates
 - It also allows a relatively simple setting for a theoretical analysis
 - But not used today; instead, medium-to-large NNs are trained using **mini-batches**

MLP Complexity

- MLPs are fast to apply but costly to train
- How can we estimate this cost?
- MLP training cost is determined by
 - Its **architecture**, that determines the number of weights to fit and that is also dependent on the input dimension D
 - The full or mini-batch **sample size** N , that determines the cost of the averages to be computed
 - The **training method**, with more or less iterations that, in turn, are more or less costly

Forward Pass Complexity

- The number of weights in a single hidden layer (SHL) MLP with D inputs, one output and one hidden layer with H units is

$$(D + 1) \times H + (H + 1) \simeq HD$$

- Each extra hidden layer pair with (H_1, H_2) units adds $(H_1 + 1) \times H_2 \simeq H_1 H_2$ weights
- The cost per pattern of the feedforward step between layers $h - 1, h$ is basically that of $H_{h-1} \times H_h$ floating point operations (FPOs)
- Thus, for a general MLP the cost in floating point operations of a forward pass is $\simeq N \times (\sum_h H_h \times H_{h-1})$
 - Very fast on GPUs as it theoretically becomes $O(N \times \sum_h H_h)$

The Cost of Computing ∇e^ℓ

- Computing a local gradient ∇e^ℓ in a SHL MLP with square error and one output essentially requires to compute
 - H components for the hidden to output connections, with a $O(1)$ cost each, for $\frac{\partial e^\ell}{\partial w_h^O} = (\hat{y} - y) o_h$
 - DH components for the input to hidden connections with essentially an $O(1)$ cost each, for $\frac{\partial e^\ell}{\partial w_{h,j}^H} = \left(\sum_{i=1}^L \delta_i w_{ih}^O \right) \sigma'(a_h) x_j$
- In these we are omitting the cost of the partials with respect the bias, but they are much smaller
- Thus, the overall cost of computing ∇e^ℓ on a SHL MLP is essentially $O(DH)$
- More layers with H_h units add a cost $O(H_{h-1} \times H_h)$
 - $O(\sum_h H_h)$ on GPUs

The Cost of Computing ∇e

- For a mini-batch of size N_b , the cost of the mini-batch gradient ∇e of a SHL MLP is $O(N_b \times D \times H)$
- And extra layers add a cost $O(N_b \times H_{h-1} \times H_h)$
- This is of the same order of magnitude than the cost of the forward pass
- This also dominates the $O(H_{h-1} \times H_h)$ cost of updating the (H_{h-1}, H_h) weights in gradient descent
- And all these costs have to be multiplied by the number of training **epochs**

Training Complexity

- The important term in the overall training cost is $n\text{Eps} \times \text{cost of } \nabla e$, with $n\text{Eps}$ the number of epochs
 - One epoch = one pass on the entire sample
- Thus, the **globally dominant term** in a SHL MLP is

$$n\text{Eps} \times \text{cost of } \nabla e = O(n\text{Eps} \times N \times D \times H)$$

with N here the entire sample size

- And for each hidden layer we add an extra cost

$$O(n\text{Eps} \times N \times H_{h-1} \times H_h)$$

- Thus, training many layered, large MLPs can be **very costly**
 - But **linear in sample size N**

Takeaways on MLP's Cost

- MLP complexity is determined by its **architecture** $\{H_h\}$, **training procedure** and **sample size** N
- The forward and backward MLP passes have basically the **same complexity**
- Their cost per **pattern and layer** is $H_{h-1} \times H_h$
- For gradient descent these costs are **multiplied by the number nEps of epochs**
- First order methods essentially **do not add extra complexity**
- Second order methods add extra per iteration costs but should **require less iterations**
- Single MLP training is **not CPU parallelizable**
- But **GPUs** can greatly improve MLP processing costs